

ML_PS1

Clarice Tee

2025-01-21

Table of contents

0.1 Data Analysis Questions	5
---------------------------------------	---

1. Familiarize w data

```
# Import required libraries
import numpy as np
import os
import statsmodels.api as sm
import pandas as pd
from sklearn.preprocessing import SplineTransformer
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
import seaborn as sns
from pygam import LinearGAM, s, te
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, r2_score, mean_squared_error
# Display all columns in pandas
pd.set_option('display.max_columns', None)
```

```
# Load the data
directory = r"C:\Users\clari\OneDrive\Documents\Machine Learning\mp1"
acs_path = os.path.join(directory, "acs_usa.csv")
acs_data = pd.read_csv(acs_path)
acs_data.describe()
print(acs_data.shape)
```

(9048, 25)

2.a. Make educdc crowsswalk

```
# Check unique values of EDUCD in the original data
acs_data['EDUCD'].unique()

array([ 65,  71,  63,  64,  81,  50, 114,  40, 101,   2,  61, 115,  26,
        30, 116,  23,  15,  25,  22,  17,  14,  16,  11,  12],
      dtype=int64)

# read in crosswalk data
crosswalk_path = os.path.join(directory, "Education-Crosswalk.csv")
crosswalk = pd.read_csv(crosswalk_path)
crosswalk.head()
```

	educd	educdc
0	2	0.0
1	10	0.0
2	11	2.0
3	12	0.0
4	13	2.5

```
# Capitalize the column names in the crosswalk
crosswalk = crosswalk.rename(columns={'educd': 'EDUCD', 'educdc': 'EDUCDC'})
# map the matching values of EDUCD between the two dataframes (i.e. merge on
↪ EDUCD)

acs_data = acs_data.merge(crosswalk, on='EDUCD', how='left')

acs_data.head()
```

	YEAR	SAMPLE	SERIAL	CBSERIAL	HHWT	CLUSTER	STRATA	GQ	PERNU
0	2023	202301	1472	2023010069599	8100.0	2023000014721	120201	4	1
1	2023	202301	2120	2023010102972	648.0	2023000021201	140401	4	1
2	2023	202301	2282	2023010110952	6642.0	2023000022821	240001	4	1
3	2023	202301	3578	2023010171383	43578.0	2023000035781	80001	4	1
4	2023	202301	3902	2023000014636	11988.0	2023000039021	150201	1	1

2.b. Making dummy variables

- i. HIGHSCHOOL 12-15 years of schooling = hs diploma

```
acs_data['HSDIP'] = np.where((acs_data['EDUCDC'] == 12) | (acs_data['EDUCDC']
↪ == 13) | (acs_data['EDUCDC'] == 14) | (acs_data['EDUCDC'] == 15),
                             1,
                             0
                             )
```

ii. COLLEGE #if more than 16 years of schooling, has college diploma daw

```
acs_data['COLDIP'] = np.where((acs_data['EDUCDC'] >= 16),
                               1,
                               0
                              )
```

iii. & iv. RACE

```
# WHITE = 1, 0 otherwise. BLACK =1, 0 otherwise
acs_data['WHITE'] = np.where(acs_data['RACE'] == 1, 1, 0)
acs_data['BLACK'] = np.where(acs_data['RACE'] == 2, 1, 0)
```

v. HISPANIC

```
acs_data['HISPANIC'] = np.where((acs_data['HISPAN'] == 1) | (acs_data['HISPAN']
↪ == 2) | (acs_data['HISPAN'] == 3) | (acs_data['HISPAN'] == 4),
                               1,
                               0
                              )
```

vi. MARRIED

```
acs_data['MAR'] = np.where(
    (acs_data['MARST'] == 1) | (acs_data['MARST'] == 2),
    1,
    0
)
```

vii. FEMALE

```
acs_data['SEX'].unique()
acs_data['FEMALE'] = np.where(acs_data['SEX'] == 2, 1, 0)
```

viii. VET (1 = VET, 0 = NOT)

```
acs_data['VET'] = np.where(acs_data['VETSTAT'] == 2, 1, 0)
```

2.c. Interaction term Create an interaction between each of the education dummy variables (hsdip and coldip) and the continuous measure of education (educdc)

```
acs_data['HSDIP_EDUCDC'] = acs_data['HSDIP'] * acs_data['EDUCDC']
acs_data['COLDIP_EDUCDC'] = acs_data['COLDIP'] * acs_data['EDUCDC']
```

2.d. Create variables

Age squared

```
# creating new variable "AGEQ" in the data, which is the square of AGE
acs_data['AGEQ'] = np.power(acs_data['AGE'], 2)
```

The natural log of incwage

```
# We first drop those with 0 income. This is because ln(0) is undefined
acs_data = acs_data[acs_data.INCWAGE != 0]
# creating new variable "INCWAGE_log" in the data, the natural log of INCWAGE
acs_data['LNINCWAGE'] = np.log(acs_data['INCWAGE'])
```

```
acs_data[['INCWAGE', 'LNINCWAGE']].head(10)
```

	INCWAGE	LNINCWAGE
0	4500	8.411833
1	5500	8.612503
2	24000	10.085809
3	26000	10.165852
4	40000	10.596635
5	70000	11.156251
6	60000	11.002100
7	30000	10.308953
8	5000	8.517193
9	118000	11.678440

0.1 Data Analysis Questions

4.1 Compute descriptive (summary) statistics for the following variables: year, incwage, lnincwage, educdc, female, age, ageq, white, black, hispanic, married, nchild, vet, hsdip, coldip, and the interaction terms. In other words, compute sample means, standard deviations, etc.

```
def acs_summary_statistics(data, variables):

    acs_summary_stats = []

    for var in variables:
        if var in data.columns: # Ensure the variable exists in the DataFrame
            stats = {
                'Variable': var,
                'Mean': data[var].mean(),
                'Std Dev': data[var].std(),
                'Min': data[var].min(),
                'Max': data[var].max(),
                'Count': data[var].count()
            }
            acs_summary_stats.append(stats)

    # Convert summary stats to a DataFrame
    return pd.DataFrame(acs_summary_stats)

# List of variables to compute summary statistics for
variables = [
    'YEAR', 'INCWAGE', 'LNINCWAGE', 'EDUCDC', 'FEMALE', 'AGE', 'AGEQ',
    'WHITE', 'BLACK', 'HISPANIC', 'MAR', 'NCHILD', 'VET',
    'HSDIP', 'COLDIP', 'HSDIP_EDUCDC', 'COLDIP_EDUCDC'
]

# Call function
acs_summary_stats = acs_summary_statistics(acs_data, variables)

# Display the results with formatting
pd.set_option('display.float_format', lambda x: '%.3f' % x)
print(acs_summary_stats.to_string(index=False))
```

Variable	Mean	Std Dev	Min	Max	Count
YEAR	2023.000	0.000	2023.000	2023.000	8610
INCWAGE	70235.557	82434.439	20.000	770000.000	8610
LNINCWAGE	10.693	1.079	2.996	13.554	8610
EDUCDC	14.310	3.054	0.000	22.000	8610
FEMALE	0.488	0.500	0.000	1.000	8610

AGE	41.627	13.291	18.000	65.000	8610
AGEQ	1909.434	1118.953	324.000	4225.000	8610
WHITE	0.663	0.473	0.000	1.000	8610
BLACK	0.082	0.275	0.000	1.000	8610
HISPANIC	0.163	0.370	0.000	1.000	8610
MAR	0.557	0.497	0.000	1.000	8610
NCHILD	0.780	1.097	0.000	6.000	8610
VET	0.041	0.198	0.000	1.000	8610
HSDIP	0.537	0.499	0.000	1.000	8610
COLDIP	0.415	0.493	0.000	1.000	8610
HSDIP_EDUCDC	6.939	6.484	0.000	14.000	8610
COLDIP_EDUCDC	7.060	8.435	0.000	22.000	8610

4.2 Scatter plot ln (incwage) and EDUCDC. Include a linear fit line. Be sure to label all axes and include an informative title.

```
# Extract data for the scatter plot
X = acs_data[['EDUCDC']].values
y = acs_data['LNINCWAGE'].values
# Fit a linear regression model
model = LinearRegression()
model.fit(X, y)
y_pred = model.predict(X)


# Create a figure and axis object
fig, ax = plt.subplots(figsize=(10, 6))

# Scatterplot of education vs. log income
ax.scatter(acs_data['EDUCDC'], acs_data['LNINCWAGE'],
           alpha=0.6, s=30, label='Data Points')

# Linear fit line
ax.plot(acs_data['EDUCDC'], y_pred, color="red",
        linewidth=2, label="Linear Fit")

# Setting the title, legend, and labels
ax.set_title(
    'Scatter Plot of Log Income (LNINCWAGE) and Education (EDUCDC)',
    ↵  fontsize=14)
ax.set_xlabel('Years of Education (EDUCDC)', fontsize=12)
ax.set_ylabel('Log Income (LNINCWAGE)', fontsize=12)
ax.legend()
ax.grid(alpha=0.3)

plt.show()
```



4.3 LNINCWAGE MODEL

```
import statsmodels.formula.api as smf

LNINCWAGE_reg = smf.ols(
    'LNINCWAGE ~ EDUCDC + FEMALE + AGE + AGEQ + WHITE + BLACK + HISPANIC+ MAR +
    ↪ NCHILD + VET', data=acs_data).fit()
print(LNINCWAGE_reg.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          LNINCWAGE    R-squared:
0.282
Model:                  OLS          Adj. R-squared:
0.281
Method:                 Least Squares    F-statistic:
337.3
Date:                   Thu, 30 Jan 2025    Prob (F-statistic):
0.00
Time:                   22:58:05          Log-Likelihood:
-11445.
No. Observations:       8610          AIC:
2.291e+04
```

Df Residuals: 8599 BIC:
2.299e+04
Df Model: 10
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025
	0.975]				
Intercept	6.1306	0.116	52.856	0.000	5.903
6.358					
EDUCDC	0.0994	0.003	29.325	0.000	0.093
0.106					
FEMALE	-0.3832	0.020	-19.100	0.000	-0.423
-0.344					
AGE	0.1493	0.006	26.016	0.000	0.138
0.161					
AGEQ	-0.0016	6.76e-05	-23.066	0.000	-0.002
-0.001					
WHITE	-0.0070	0.028	-0.253	0.801	-0.061
0.047					
BLACK	-0.1202	0.043	-2.791	0.005	-0.205
-0.036					
HISPANIC	-0.0211	0.033	-0.640	0.522	-0.086
0.043					
MAR	0.2053	0.023	8.867	0.000	0.160
0.251					
NCHILD	-0.0082	0.010	-0.794	0.427	-0.028
0.012					
VET	-0.0132	0.050	-0.261	0.794	-0.112
0.086					
Omnibus:	2566.758		Durbin-Watson:		
1.845					
Prob(Omnibus):	0.000		Jarque-Bera (JB):		
11879.119					
Skew:	-1.379		Prob(JB):		
0.00					
Kurtosis:	8.050		Cond. No.		
2.62e+04					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.62e+04. This might indicate that there are strong multicollinearity or other numerical problems.

4.3.a The model explains 28.1% of the variation in log wages (adjusted r2 value)

4.3.b. (b) What is the return to an additional year of education? Is this statistically significant? Is it practically significant? Briefly explain. note that your answer should be in terms of wages, not log wages.

```
beta_EDUCDC = LNINCWAGE_reg.params.iloc[1]

# Convert
change_wages_EDUCDC = (np.exp(beta_EDUCDC) - 1) * 100

mean_wage = acs_data['INCWAGE'].mean() # Calculate mean directly from the data

# Calculate dollar change
dollar_change_EDUCDC = mean_wage * (change_wages_EDUCDC / 100)

print("Additional year of education means a 10.45 percent increase in wages.
↪ Dollar change in wages for an additional year of education is:",
↪ dollar_change_EDUCDC)
```

Additional year of education means a 10.45 percent increase in wages.
Dollar change in wages for an additional year of education is:
7341.363608190367

Based on the p-value, it is statistically significant at the 1% level.

4.3.c.

Here, we need to take the derivative of the model w rt AGE and equate it to 0. We are left with $\text{Beta}(3) + 2\text{Beta}(4)\text{AGE}$

```
# Getting the coefficients for AGE and AGE^2
beta_AGE = LNINCWAGE_reg.params.iloc[3]
beta_AGE2 = LNINCWAGE_reg.params.iloc[4]

# Calculate the age at which the highest wage occurs
max_age = -beta_AGE / (2 * beta_AGE2)
print(max_age)
```

47.861998868799084

We must check whether this is a local max or minimum by taking the second derivative wrt age, leaving us with $2\beta_4$.

```
second_deriv = (2 * beta_AGE2)
print(second_deriv)
```

-0.0031195396677629757

Now that we know it is negative, we know this is the maximum. therefore:

```
print(f"The age at which the model predicts the highest wage is approximately  
↪ 48 years.")
```

The age at which the model predicts the highest wage is approximately 48 years.

4.3.d.

```
# Getting the Coefficient for FEMALE
beta_female = LNINCWAGE_reg.params.iloc[2]

# Convert log-wage to percentage change
percentage_change_female = (np.exp(beta_female) - 1) * 100

print(f"Being female is associated with a {percentage_change_female:.1f}%  
↪ change in wages.")
```

Being female is associated with a -31.8% change in wages.

The model predicts that being a woman means that, all else equal, wages are lower by about 31.8%. This makes sense because of gender discrimination against women, where their work is undervalued relative to men or they may take lower paying jobs due to having to work part-time to take on child rearing work.

4.3.e. All else equal, white individuals have approximately a .7% decrease in their wages compared to other races, although, this is not significant because we cannot even reject the null hypothesis at a significance level of 10%.

All else equal, black individuals have 12.02% less than other races and this proves to be a significant value at the 1% level.

4.4

```

## Create degree groups
acs_data['GROUP'] = 'No High School Diploma' # Default group
acs_data.loc[(acs_data['HSDIP'] == 1) & (acs_data['COLDIP'] == 0), 'GROUP'] =
    ↪ 'High School Diploma'
acs_data.loc[(acs_data['COLDIP'] == 1), 'GROUP'] = 'College Degree'

# Create the plot
plt.figure(figsize=(12, 8))

# Colors per grp
group_colors = {
    'No High School Diploma': 'orange',
    'High School Diploma': 'green',
    'College Degree': 'blue'
}

# Plot scatter points and regression lines for each group
for group, color in group_colors.items():
    subset = acs_data[acs_data['GROUP'] == group] # Filter data for each group

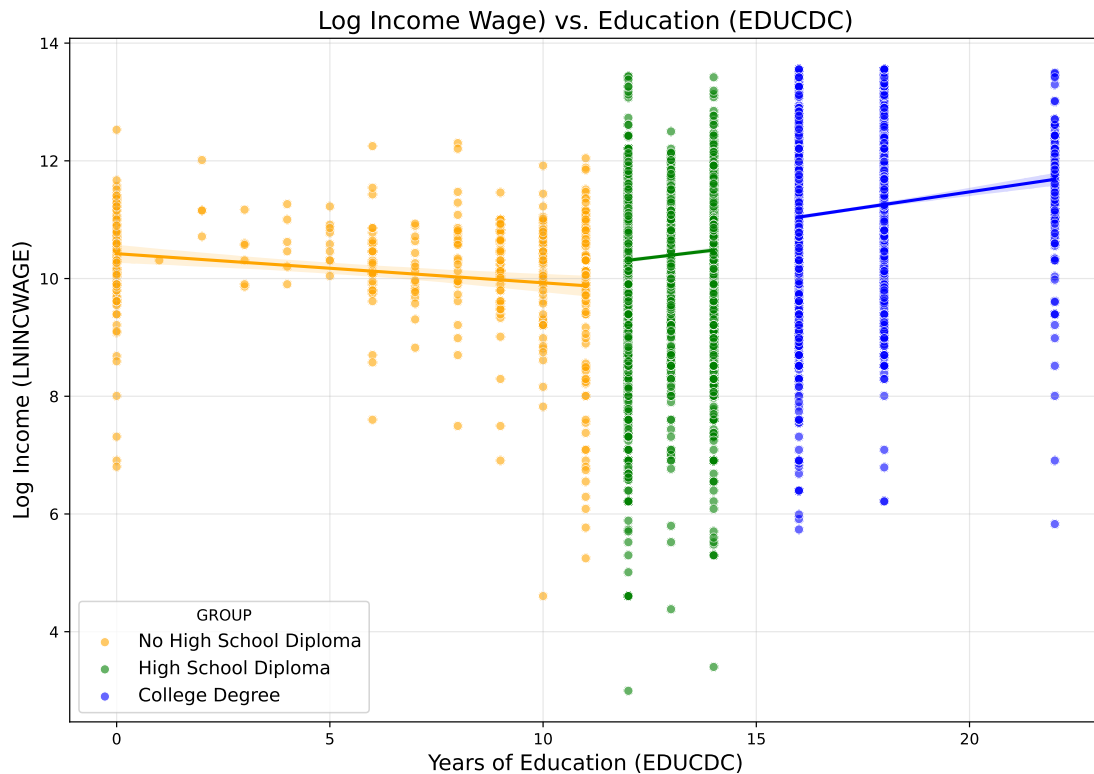
    # Scatterplot for this group
    sns.scatterplot(
        x='EDUCDC',
        y='LNINCWAGE',
        data=subset,
        color=color,
        alpha=0.6,
        label=group
    )

    # Regression line for this group
    sns.regplot(
        x='EDUCDC',
        y='LNINCWAGE',
        data=subset,
        scatter=False,
        color=color,
        line_kws={'linewidth': 2},
        truncate=True
    )

# Add labels, title, and legend
plt.title('Log Income Wage) vs. Education (EDUCDC)', fontsize=16)
plt.xlabel('Years of Education (EDUCDC)', fontsize=14)
plt.ylabel('Log Income (LNINCWAGE)', fontsize=14) # Fix label
plt.legend(title='GROUP', fontsize=12)
plt.grid(alpha=0.3)

```

```
# Show the plot
plt.show()
```



4.5a $LNINCWAGES = \beta_0 + \beta_1(hsdip) + \beta_2(coldip) + \beta_3(educdc \times hsdip) + \beta_4(educdc \times coldip) + \beta_5(age) + \beta_6(age^2) + \beta_7(female) + \beta_8(white) + \beta_9(black) + \beta_{10}(hispanic) + \beta_{11}(married) + \beta_{12}(nchild) + \beta_{13}(vet) + \epsilon$

Using theory, intuition, and/or common sense, explain/justify why you think your model is the best possible representation of the way the world works (in other words, why you think you are correctly modeling $f(X)$ but not over-fitting e).

By now including the dummy variables for HSDIP, COLDIP and their interaction terms with EDUCDC, the regression model can show differential returns to education based on educational attainment. We increase flexibility of the model. Now, there can be different slopes for the relationship between educational attainment and wage. It shows us potential “sheepskin effects”, so we can see the added value of completing a HS or College—the degrees have a signalling effect that help employers quickly screen competence. https://www.econlib.org/archives/2012/01/the_present_val.html

4.5b

```
NEW_reg = smf.ols(
    'LNINCWAGE ~ HSDIP + COLDIP + HSDIP_EDUCDC + COLDIP_EDUCDC + FEMALE + AGE +
    ↪ AGEQ + WHITE + BLACK + HISPANIC+ MAR + NCHILD + VET',
    ↪ data=acs_data).fit()
print(NEW_reg.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          LNINCWAGE    R-squared:
0.306
Model:                  OLS          Adj. R-squared:
0.305
Method:                 Least Squares    F-statistic:
292.0
Date:                   Thu, 30 Jan 2025    Prob (F-statistic):
0.00
Time:                   22:58:09          Log-Likelihood:
-11296.
No. Observations:      8610          AIC:
2.262e+04
Df Residuals:          8596          BIC:
2.272e+04
Df Model:              13
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025
	0.975]				

Intercept	7.1693	0.118	60.790	0.000	6.938
7.401					
HSDIP	-0.7122	0.189	-3.761	0.000	-1.083
-0.341					
COLDIP	-0.1929	0.177	-1.090	0.276	-0.540
0.154					
HSDIP_EDUCDC	0.0833	0.014	5.855	0.000	0.055
0.111					
COLDIP_EDUCDC	0.0708	0.010	7.063	0.000	0.051
0.090					
FEMALE	-0.3967	0.020	-20.056	0.000	-0.435
-0.358					
AGE	0.1368	0.006	24.069	0.000	0.126
0.148					

AGEQ	-0.0014	6.7e-05	-21.149	0.000	-0.002
-0.001					
WHITE	0.0238	0.027	0.875	0.382	-0.030
0.077					
BLACK	-0.0745	0.042	-1.756	0.079	-0.158
0.009					
HISPANIC	-0.0090	0.032	-0.278	0.781	-0.073
0.055					
MAR	0.1833	0.023	8.039	0.000	0.139
0.228					
NCHILD	-0.0029	0.010	-0.290	0.772	-0.023
0.017					
VET	0.0199	0.050	0.399	0.690	-0.078
0.117					

```
=====
Omnibus:                2730.221   Durbin-Watson:
1.856
Prob(Omnibus):           0.000   Jarque-Bera (JB):
13032.061
Skew:                    -1.465   Prob(JB):
0.00
Kurtosis:                8.267   Cond. No.
4.41e+04
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.41e+04. This might indicate that there are strong multicollinearity or other numerical problems.

REsults: The following variables are statistically significant at the 10% level: HSDIP, HSDIP_EDUCDC, COLDIP_EDUCDC, FEMALE, AGE, AGEQ, BLACK.

Importantly, our interaction terms proved to be statistically significant, confirming the sheepskin effect or the idea that an HS degree and a college degree bumps up wages.

4.5c

```
# 22 year old, female, not white, black or hispanic. Not married, no children,
↪ not a veteran, HSDIP.
# vs all else equal, but COLDIP
female_1 = acs_data.loc[
    (acs_data['AGE'] == 22) & (acs_data['FEMALE'] == 1) & (acs_data['WHITE'] ==
↪ 0) & (
```

```

        acs_data['BLACK'] == 0) & (acs_data['NCHILD'] == 0) & (acs_data['VET']
↪ == 0) & (acs_data['MAR'] == 0) & (acs_data['HSDIP'] == 1) &
↪ (acs_data['COLDIP'] == 0) & (acs_data['HSDIP_EDUCDC'] == 12)&
↪ (acs_data['COLDIP_EDUCDC'] == 0)
    ]
    # Expectation
    female_1_prediction = LNINCWAGE_reg.get_prediction(female_1)
    # Display prediction summary
    prediction_summary = female_1_prediction.summary_frame(alpha=0.05)
    print(prediction_summary)

    # Convert log wages to actual wages
    predicted_wage = np.exp(prediction_summary['mean'][0])
    print(f"Predicted wage with a high school diploma: ${predicted_wage:.2f}")

    # alpha : significance level for the confidence interval.
    # Fix alpha to 0.05 so that the confidence interval should have 95% coverage.

```

	mean	mean_se	mean_ci_lower	mean_ci_upper	obs_ci_lower	obs_ci_upper
0	9.470	0.035	9.402	9.539	7.675	11.265
1	9.470	0.035	9.402	9.539	7.675	11.265
2	9.449	0.033	9.385	9.513	7.655	11.244
3	9.449	0.033	9.385	9.513	7.655	11.244
4	9.470	0.035	9.402	9.539	7.675	11.265
5	9.449	0.033	9.385	9.513	7.655	11.244
6	9.449	0.033	9.385	9.513	7.655	11.244
7	9.449	0.033	9.385	9.513	7.655	11.244
8	9.470	0.035	9.402	9.539	7.675	11.265
9	9.449	0.033	9.385	9.513	7.655	11.244
10	9.449	0.033	9.385	9.513	7.655	11.244
11	9.449	0.033	9.385	9.513	7.655	11.244

12 9.470 0.035 9.402 9.539 7.675
11.265

Predicted wage with a high school diploma: \$12966.88

COLDIP

```
# 22 year old, female, not white, black or hispanic. Not married, no children,
↪ not a veteran, HSDIP.
# vs all else equal, but COLDIP
female_2 = acs_data.loc[
    (acs_data['AGE'] == 22) & (acs_data['FEMALE'] == 1) & (acs_data['WHITE'] ==
↪ 0) & (
        acs_data['BLACK'] == 0) & (acs_data['NCHILD'] == 0) & (acs_data['VET']
↪ == 0) & (acs_data['MAR'] == 0) & (acs_data['HSDIP'] == 0) &
↪ (acs_data['COLDIP'] == 1) & (acs_data['HSDIP_EDUCDC'] == 0)&
↪ (acs_data['COLDIP_EDUCDC'] == 16)
]
# Expectation
female_2_prediction = LNINCWAGE_reg.get_prediction(female_2)
# Display prediction summary
prediction_summary_2 = female_2_prediction.summary_frame(alpha=0.05)
print(prediction_summary_2)

# Convert log wages to actual wages
predicted_wage_2 = np.exp(prediction_summary_2['mean'][0])
print(f"Predicted wage with a college diploma: ${predicted_wage_2:.2f}")

# alpha : significance level for the confidence interval.
# Fix alpha to 0.05 so that the confidence interval should have 95% coverage.
```

	mean	mean_se	mean_ci_lower	mean_ci_upper	obs_ci_lower
0	9.868	0.035	9.800	9.936	8.073
11.662					
1	9.868	0.035	9.800	9.936	8.073
11.662					
2	9.847	0.035	9.779	9.915	8.052
11.641					
3	9.868	0.035	9.800	9.936	8.073
11.662					
4	9.847	0.035	9.779	9.915	8.052
11.641					
5	9.847	0.035	9.779	9.915	8.052
11.641					
6	9.868	0.035	9.800	9.936	8.073
11.662					

7	9.868	0.035	9.800	9.936	8.073
	11.662				
8	9.868	0.035	9.800	9.936	8.073
	11.662				
9	9.868	0.035	9.800	9.936	8.073
	11.662				

Predicted wage with a college diploma: \$19299.12

4.5d Those with college degrees have a higher predicted wage than those without. They earn about 6, 332 USD more on average.

4.5e Since we know that a college degree will increase the average expected wage, it would be good to create policy to expand access to college education. However, what our model does is to test correlation and not causation, so we do need to consider how other factors influence this (omitted variables like household size or grades in school). The legislation should also try to ensure access to good quality education and ensure that the quality and type of education matter, not just access. And we should consider who (race, gender, origin, socioeconomic status) is being provided with this subsidy.

4.5f

```
NEW_reg_r_squared = NEW_reg.rsquared
NEW_reg_adjusted_r_squared = round(NEW_reg.rsquared_adj,4)

print(f"R-squared: {NEW_reg_r_squared:.4f}")
print(f"{{(NEW_reg_adjusted_r_squared*100)}} percent of the variation in logwages
↪ can be explained by the model")
```

R-squared: 0.3063

30.53 percent of the variation in logwages can be explained by the model

versus

```
OLD_reg_r_squared = LNINCWAGE_reg.rsquared
OLD_reg_adjusted_r_squared = LNINCWAGE_reg.rsquared_adj

print(f"R-squared: {OLD_reg_r_squared:.4f}")
print(f"Adjusted R-squared: {OLD_reg_adjusted_r_squared:.4f}")
```

R-squared: 0.2817

Adjusted R-squared: 0.2809

We see that the new model has a higher R² and adjusted R², meaning it better explains the variation in log wages compared to the model from 3.

4.5g

For both 22 year olds, the CIs are pretty narrow, meaning the estimates are likely quite precise. High school diploma, the model predicts a log wage of 9.470154 with a 95% confidence interval of [9.401522, 9.538785]. For the college graduate, the predicted log wage is 9.867815 with a confidence interval of [9.799870, 9.935759]. But, the model only explains 30.53% of the variation in log wages, so there's still quite a bit of variation that is not explained.

So, while our model is quite precise, we must be wary. We may need to explore a bit more and use data from previous years as well.

4.6a

```
# Predictors from Question 3 (excluding AGE and AGE_SQUARED)
predictors_q3 = ["EDUCDC", "FEMALE", "WHITE", "BLACK", "HISPANIC",
                 "MAR", "NCHILD", "VET"]

# Define the age variable and the target variable
X_age = acs_data[["AGE"]]
y = acs_data["LNINCWAGE"]

# Define the knots for the spline (18 and 65)
knots = np.array([18, 65]).reshape(-1, 1)

# Create the spline transformer
spline_transformer = SplineTransformer(degree=3, knots=knots,
                                       include_bias=False)

# Transform the age variable using the spline
X_splines = spline_transformer.fit_transform(X_age)

# Combine the spline-transformed age variable with the other predictors
X_controls = acs_data[predictors_q3]
X = np.hstack([X_splines, X_controls])

# Fit the linear regression model
model = LinearRegression()
model.fit(X, y)

# Print the intercept and coefficients
print("Intercept:", model.intercept_)
spline_columns = [f"spline_{i}" for i in range(X_splines.shape[1])]
all_columns = spline_columns + predictors_q3
coefficients = pd.Series(model.coef_, index=all_columns)
print("Coefficients:")
print(coefficients)

# Calculate the adjusted R-squared
r_squared = model.score(X, y)
```

```

n = len(y)
p = X.shape[1]
adjusted_r_squared = 1 - ((1 - r_squared) * (n - 1)) / (n - p - 1)
print("The Adjusted R-squared of the model with the two age knots is",
      ↪ adjusted_r_squared)

```

Intercept: 14.075815645865092

Coefficients:

```

spline_0    -22.189
spline_1     -1.931
spline_2     -6.236
EDUCDC       0.096
FEMALE       -0.379
WHITE        -0.013
BLACK        -0.130
HISPANIC     -0.028
MAR          0.190
NCHILD       -0.008
VET          -0.016

```

dtype: float64

The Adjusted R-squared of the model with the two age knots is
0.2914488777671098

4.6b The adjusted R² of .291 is slightly higher than the one in model 3. The adjusted R² of .291 is higher than the one in model 3. This is probably because, as we learned, splines capture nonlinear relationships. In this case, the nonlinear effects of age, which thus improves the model's fit (it explains more of the variation).

4.6c

```

# Define the knots for the spline (24 and 55)
knots_24_55 = np.array([24, 55]).reshape(-1, 1)

# Create the spline transformer
spline_transformer_24_55 = SplineTransformer(degree=3, knots=knots_24_55,
      ↪ include_bias=False)

# Transform the age variable using the spline
X_splines_24_55 = spline_transformer_24_55.fit_transform(X_age)

# Combine the spline-transformed age variable with the other predictors
X_24_55 = np.hstack([X_splines_24_55, X_controls])

# Fit the linear regression model
model_24_55 = LinearRegression()
model_24_55.fit(X_24_55, y)

```

```

# Print the intercept and coefficients
print("Intercept (knots at 24 and 55):", model_24_55.intercept_)
spline_columns_24_55 = [f"spline_24_55_{i}" for i in
    ↪ range(X_splines_24_55.shape[1])]
all_columns_24_55 = spline_columns_24_55 + predictors_q3
coefficients_24_55 = pd.Series(model_24_55.coef_, index=all_columns_24_55)
print("Coefficients (knots at 24 and 55):")
print(coefficients_24_55)

# Calculate the adjusted R-squared
r_squared_24_55 = model_24_55.score(X_24_55, y)
n_24_55 = len(y)
p_24_55 = X_24_55.shape[1]
adjusted_r_squared_24_55 = 1 - ((1 - r_squared_24_55) * (n_24_55 - 1)) /
    ↪ (n_24_55 - p_24_55 - 1)
print("Adjusted R-squared (knots at 24 and 55):", adjusted_r_squared_24_55)

```

```

Intercept (knots at 24 and 55): 12.615808625187793
Coefficients (knots at 24 and 55):
spline_24_55_0    -13.896
spline_24_55_1     -1.510
spline_24_55_2     -4.272
EDUCDC              0.102
FEMALE              -0.385
WHITE               -0.015
BLACK               -0.127
HISPANIC            -0.016
MAR                  0.211
NCHILD              0.008
VET                 -0.019
dtype: float64
Adjusted R-squared (knots at 24 and 55): 0.2673685975131527

```

Model of my choosing

```

# Knots for the spline (21 and 65)
knots_21_65 = np.array([21, 65]).reshape(-1, 1)

# Create the spline transformer
spline_transformer_21_65 = SplineTransformer(degree=3, knots=knots_21_65,
    ↪ include_bias=False)

# Transform the age variable using the spline
X_splines_21_65 = spline_transformer_21_65.fit_transform(X_age)

```

```

# Combine the spline-transformed age variable with the other predictors
X_21_65 = np.hstack([X_splines_21_65, X_controls])

# Fit the linear regression model
model_21_65 = LinearRegression()
model_21_65.fit(X_21_65, y)

# Print the intercept and coefficients
print("Intercept (knots at 21 and 65):", model_21_65.intercept_)
spline_columns_21_65 = [f"spline_21_65_{i}" for i in
    ↪ range(X_splines_21_65.shape[1])]
all_columns_21_65 = spline_columns_21_65 + predictors_q3
coefficients_21_65 = pd.Series(model_21_65.coef_, index=all_columns_21_65)
print("Coefficients (knots at 21 and 65):")
print(coefficients_21_65)

# Getting the adjusted R-squared
r_squared_21_65 = model_21_65.score(X_21_65, y)
n_21_65 = len(y)
p_21_65 = X_21_65.shape[1]
adjusted_r_squared_21_65 = 1 - ((1 - r_squared_21_65) * (n_21_65 - 1)) /
    ↪ (n_21_65 - p_21_65 - 1)
print("Adjusted R-squared (knots at 21 and 65):", adjusted_r_squared_21_65)

```

```

Intercept (knots at 21 and 65): 13.613217907092636
Coefficients (knots at 21 and 65):
spline_21_65_0    -18.378
spline_21_65_1     -1.803
spline_21_65_2    -5.629
EDUCDC             0.098
FEMALE             -0.380
WHITE              -0.013
BLACK              -0.127
HISPANIC           -0.022
MAR                0.198
NCHILD            -0.006
VET                -0.014
dtype: float64
Adjusted R-squared (knots at 21 and 65): 0.2819850168684307

```

Since my model of 21 and 65 have a higher adjusted R², it is a more powerful model that fits the data better than the suggested one. I chose 21, thinking this is the age in which many students graduate from university and 65 is when they retire, which would likely improve the model.

4.6d

```

# Create a test dataframe for predictions
df_test = pd.DataFrame({
    "EDUCDC": [16, 16],
    "FEMALE": [1, 1],
    "AGE": [17, 50],
    "WHITE": [0, 0],
    "BLACK": [0, 0],
    "HISPANIC": [0, 0],
    "MAR": [0, 0],
    "NCHILD": [0, 0],
    "VET": [0, 0]
})

# Transform AGE into splines using the transformer from 6.c (knots at 24 and
↪ 55)
splines_test = spline_transformer_24_55.transform(df_test[["AGE"]])

# Combine spline-transformed AGE with other predictors from Question 3
X_test = np.hstack([splines_test, df_test[predictors_q3]])

# Predict log wages using the model with knots at 24 and 55
log_wage_pred = model_24_55.predict(X_test)

# Convert log wages to actual wages
wage_pred = np.exp(log_wage_pred)

# Create a results dataframe
df_pred = pd.DataFrame({
    "AGE": [17, 50],
    "Predicted Log Wage": log_wage_pred,
    "Predicted Wage": wage_pred
})

print(df_pred)

```

	AGE	Predicted Log Wage	Predicted Wage
0	17	9.824	18470.276
1	50	10.713	44950.928

The knots in the spline let us reflect wage differences because of the non-linear effects of age.

Here, we are assuming both the 17 and 50 year olds have 16 years of education, even if it isn't really realistic. Predictably, the 17 year olds, who have less work experience have lower predicted wage by around 31K USD. Meanwhile, at 50, one is likely at the top—a managing partner, director, VP, department head, etc, making earnings much higher.