

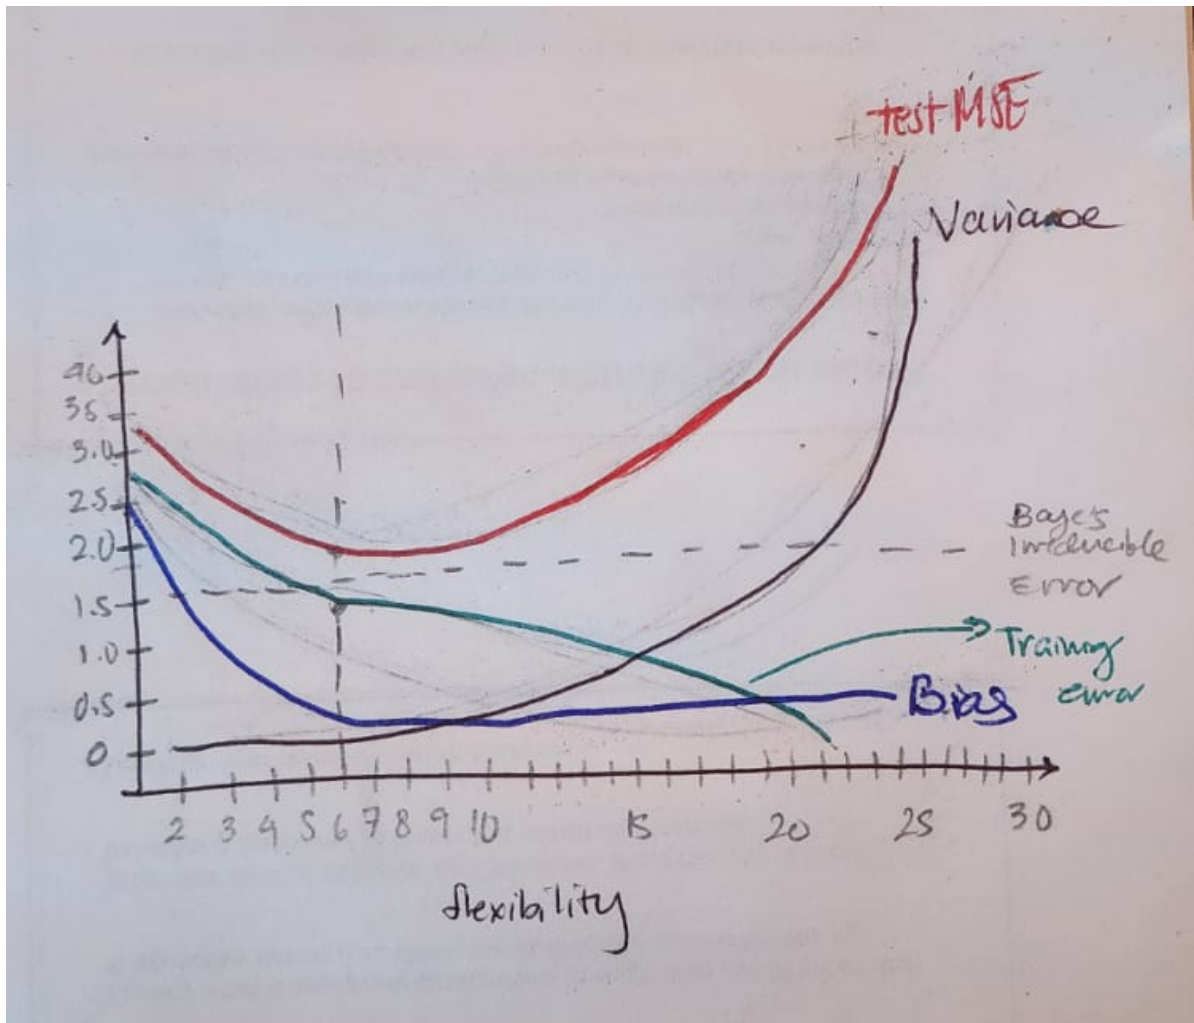
ML_PS1

Clarice Tee

2025-01-21

Part 1

```
picture =r'C:\Users\clari\OneDrive\Documents\Machine Learning\ps1'
```



1.a.

1.b. squared bias-this goes down monotonically until it flattens out because greater flexibility leads to the model better being able to capture the true relationship in the data, until a certain point (diminishing returns) where it doesn't add much else.

irreducible error- this is the lower limit of the test MSE, thus it is a straight line. The test MSE is above this line and the part of the training MSE below this line is when the data has been overfitted.

test MSE- its shape is concave, with the curve going upwards because more flexibility yields a better fit, until it overfits.

training MSE- this error goes down monotonically since more flexibility leads to better fitting data.

variance- this increases monotonically as flexibility increases, up to a point where there is overfitting.

2. Advantages (flexible): You have less bias and a better fit for non-linear models.

Disdvantages (flexible): Because we are estimating more parameters, it's more likely that we have overfitting from having too much noise. THis also means greater variance observed in the model.

A more flexible approach would be better if we want to capture more complex patterns from a larger sample size. This is suitable when we care more about the prediction we can make with the data, although there may be more variance.

A less flexible is better when we have a less data and are more interested in trying to interpret or make sense of our results. However, we are more likely to have biased restuls.

3.a.

```
# Import the libraries
import numpy as np
import os
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.formula.api as smf
from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, r2_score, mean_squared_error
```

3.b. There are 506 rows and 14 columns. The rows are the census tracts in Boston and the columns are the predictors. Based on the text file, the variables reporesent these things specifically:

```
# Load the dataset
directory = r"C:\Users\clari\OneDrive\Documents\Machine
↳ Learning\ps1\Data-Boston\Boston"
boston_path = os.path.join(directory, "Boston.csv")
boston_df = pd.read_csv(boston_path)
print(boston_df.dtypes)
print(boston_df.shape)
```

CRIM	float64
ZN	float64
INDUS	float64
CHAS	float64

```

NOX          float64
RM           float64
AGE          float64
DIS          float64
RAD          float64
TAX          float64
PTRATIO      float64
B            float64
LSTAT        float64
MDEV         float64
dtype: object
(506, 14)

```

CRIM per capita crime rate by town ZN proportion of residential land zoned for lots over 25,000 sq.ft. INDUS proportion of non-retail business acres per town CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise) NOX nitric oxides concentration (parts per 10 million) RM average number of rooms per dwelling AGE proportion of owner-occupied units built prior to 1940 DIS weighted distances to five Boston employment centres RAD index of accessibility to radial highways TAX full-value property-tax rate per \$10,000 PTRATIO pupil-teacher ratio by town B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town LSTAT % lower status of the population MEDV Median value of owner-occupied homes in \$1000's

```

# Lookup the missing values
print(boston_df.isna().sum())

```

```

CRIM      0
ZN         0
INDUS      0
CHAS       0
NOX        0
RM         0
AGE        0
DIS        0
RAD        0
TAX        0
PTRATIO    0
B          0
LSTAT      0
MDEV       0
dtype: int64

```

3.c.

```
pairwise = sns.pairplot(boston_df, vars = ['AGE', 'CRIM', 'MDEV', 'B', 'DIS'])  
plt.show()
```

C:\Users\clari\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:

FutureWarning:

use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.

C:\Users\clari\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:

FutureWarning:

use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.

C:\Users\clari\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:

FutureWarning:

use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.

C:\Users\clari\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:

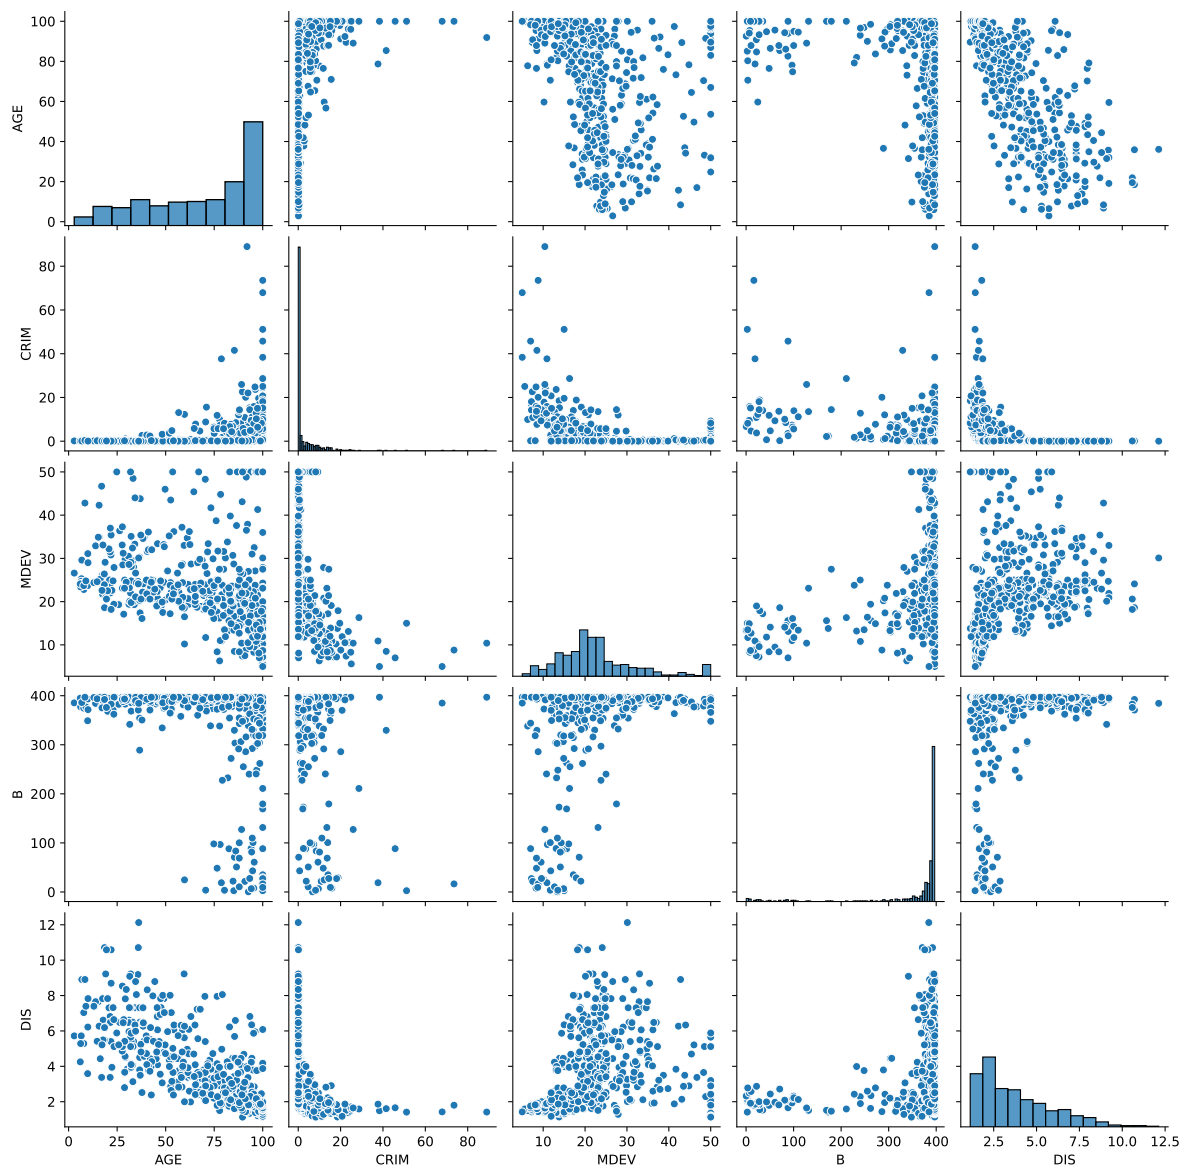
FutureWarning:

use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.

C:\Users\clari\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:

FutureWarning:

use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.



Age vs CRIM: More old buildings, more observations of crime and higher per capita crime rates by town

B vs DIS: Towns with a higher proportion of Black people are further from the Boston employment centers.

MDEV vs CRIM In homes not occupied by owners, more observations of crime and higher per capita crime rates by town

DIS vs CRIM: There seems to be more and higher per capita crime rates by town when closer to the Boston employment centers.

DIS vs AGE: It looks like a negative linear relationship between the two. Towns with more pre 1940's buildings are closer to the Boston employment centers.

3.d. All the others that I chose seem to have a somewhat similar pattern. AGE and B, it is more like a positive correlation, meaning there are more observations and higher per capita crime rates in places with a higher proportion of older homes and Black population. For MDEV and DIS, it looks like a negative correlation. However, none of these look like a clear linear relationship with crime. The data is even highly varied at some points.


3.e.

```
#Box plot columns
box_values = ["CRIM", "TAX", "PTRATIO"]

fig, axs = plt.subplots(1, 3, figsize=(16, 9))

for i, col in enumerate(box_values):
    sns.boxplot(y=boston_df[col], ax=axs[i], color='purple')
    axs[i].set_title(f"{col}")
    axs[i].set_ylabel(col)
    axs[i].set_xlabel("")

plt.tight_layout()
plt.show()
```



Reference Lab and <https://stackoverflow.com/questions/48176920/how-to-iteratively-plot-different-data-as-boxplots-in-seaborn-without-them-over>

CRIM has the largest range. There is a small interquartiel range (most cities have low per capita crime rates) and a lot of outliers (values that fall outside of the min max whixkers). This may indicate some data measurement or encoding errors or something else that is happening that must be investigated.

TAX has the narrowest range of the three. It seems that the broad IQR caputes most of the data, leaving no outliers.

Pupil-teacher ratios have a small range, the highest median among the three and have a large whixker range relative to the IQR(the min and max values are far from the middle values) and there are outliers. It is also more skewed to higher ratios

3.f. There are 35 census tracts that bound the Charles River.

```
chas_subset = boston_df[boston_df['CHAS'] == 1]
chas_count = chas_subset.shape[0]

print(f"There are {chas_count} tracts that bound the Charles River")
```

There are 35 tracts that bound the Charles River

3.g.


```
md_ptratio = boston_df["PTRATIO"].median()

print(f"The median PTRATIO is {md_ptratio}.")
```

The median PTRATIO is 19.05.

3.h. Get the min MDEV

```
boston_df_sort = boston_df.sort_values("MDEV")

min_MDEV = boston_df_sort.head(5)

# Print the row
print(min_MDEV)
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
398	38.35180	0.0	18.1	0.0	0.693	5.453	100.0	1.4896	24.0	666.0	
405	67.92080	0.0	18.1	0.0	0.693	5.683	100.0	1.4254	24.0	666.0	
400	25.04610	0.0	18.1	0.0	0.693	5.987	100.0	1.5888	24.0	666.0	
399	9.91655	0.0	18.1	0.0	0.693	5.852	77.8	1.5004	24.0	666.0	
414	45.74610	0.0	18.1	0.0	0.693	4.519	100.0	1.6582	24.0	666.0	

	PTRATIO	B	LSTAT	MDEV
398	20.2	396.90	30.59	5.0
405	20.2	384.97	22.98	5.0
400	20.2	396.90	26.77	5.6
399	20.2	338.16	29.97	6.3
414	20.2	88.27	36.98	7.0

Tracts 398 and 405 have the minimum MDEV value of 5.

```
min_MDEV_rows = boston_df.loc[[398, 405]]

# Getting the summary statistics for all predictors
summary_stats_MDEV = boston_df.describe().T[['min', 'max', 'mean']]

# Add the columns w the values from 398 and 405 to the summary
summary_stats_MDEV['Row_398'] = min_MDEV_rows.iloc[0]
summary_stats_MDEV['Row_405'] = min_MDEV_rows.iloc[1]

print(summary_stats_MDEV)
```

	min	max	mean	Row_398	Row_405
CRIM	0.00632	88.9762	3.593761	38.3518	67.9208
ZN	0.00000	100.0000	11.363636	0.0000	0.0000
INDUS	0.46000	27.7400	11.136779	18.1000	18.1000
CHAS	0.00000	1.0000	0.069170	0.0000	0.0000
NOX	0.38500	0.8710	0.554695	0.6930	0.6930
RM	3.56100	8.7800	6.284634	5.4530	5.6830
AGE	2.90000	100.0000	68.574901	100.0000	100.0000
DIS	1.12960	12.1265	3.795043	1.4896	1.4254
RAD	1.00000	24.0000	9.549407	24.0000	24.0000
TAX	187.00000	711.0000	408.237154	666.0000	666.0000
PTRATIO	12.60000	22.0000	18.455534	20.2000	20.2000
B	0.32000	396.9000	356.674032	396.9000	384.9700
LSTAT	1.73000	37.9700	12.653063	30.5900	22.9800
MDEV	5.00000	50.0000	22.532806	5.0000	5.0000

CRIM: Areas with this MDEV have per capita crime rates far above the min and mean. They are outlier values (see boxplot)

AGE: Areas with this MDEV are outliers, with their AGE values being the maximum.

DIS: Areas with this MDEV fall within the range of values, below the mean and above the minimum.

RAD: Areas with this MDEV are outliers, with their RAD values being the maximum.

CHAS: They are not bounded by the Charles river

TAX: Areas with this MDEV lie outside of the IQR (see boxplot), with their TAX values being closer to the maximum.

PTRATIO: Areas with this MDEV fall within the range, but are higher than the mean, closer to the max value.

LSTAT: Areas with this MDEV fall within the range, but are higher than the mean, closer to the max value.

3.i.

```
# RM > 7
RM_7 = boston_df[boston_df["RM"] > 7]
RM_7_count = RM_7.shape[0]

# RM > 8
RM_8 = boston_df[boston_df["RM"] > 8]
RM_8_count = RM_8.shape[0]
```

```

print(f"There are {RM_7_count} rooms with an average of more than 7 rooms per
↪ dwelling")

print(f"There are {RM_8_count} rooms with an average of more than 8 rooms per
↪ dwelling")

```

There are 64 rooms with an average of more than 7 rooms per dwelling
There are 13 rooms with an average of more than 8 rooms per dwelling

Getting the summary stats

```

# Summary stats for RM > 7
RM_8_summary = RM_8.agg(['mean', 'min', 'max']).T

# Compute overall mean, min, and max for comparison
summary_stats_RM = boston_df.describe().T[['min', 'max', 'mean']]

# Combine the two summaries
comparison_summary_RM_7 = summary_stats_RM.join(RM_8_summary,
↪ lsuffix='_Overall', rsuffix='_RM>8')

# Display the comparison
print(comparison_summary_RM_7)

# Summary stats for RM > 8
RM_8_summary = RM_8.agg(['mean', 'min', 'max']).T

# Compute overall mean, min, and max for comparison
summary_stats_RM = boston_df.describe().T[['min', 'max', 'mean']]

# Combine the two summaries
comparison_summary_RM_8 = summary_stats_RM.join(RM_8_summary,
↪ lsuffix='_Overall', rsuffix='_RM>8')

# Display the comparison
print(comparison_summary_RM_8)

```

	min_Overall	max_Overall	mean_Overall	mean_RM>8	min_RM>8 \
CRIM	0.00632	88.9762	3.593761	0.718795	0.02009
ZN	0.00000	100.0000	11.363636	13.615385	0.00000
INDUS	0.46000	27.7400	11.136779	7.078462	2.68000

CHAS	0.00000	1.0000	0.069170	0.153846	0.00000
NOX	0.38500	0.8710	0.554695	0.539238	0.41610
RM	3.56100	8.7800	6.284634	8.348538	8.03400
AGE	2.90000	100.0000	68.574901	71.538462	8.40000
DIS	1.12960	12.1265	3.795043	3.430192	1.80100
RAD	1.00000	24.0000	9.549407	7.461538	2.00000
TAX	187.00000	711.0000	408.237154	325.076923	224.00000
PTRATIO	12.60000	22.0000	18.455534	16.361538	13.00000
B	0.32000	396.9000	356.674032	385.210769	354.55000
LSTAT	1.73000	37.9700	12.653063	4.310000	2.47000
MDEV	5.00000	50.0000	22.532806	44.200000	21.90000

	max_RM>8
CRIM	3.47428
ZN	95.00000
INDUS	19.58000
CHAS	1.00000
NOX	0.71800
RM	8.78000
AGE	93.90000
DIS	8.90670
RAD	24.00000
TAX	666.00000
PTRATIO	20.20000
B	396.90000
LSTAT	7.44000
MDEV	50.00000

	min_Overall	max_Overall	mean_Overall	mean_RM>8	min_RM>8 \
CRIM	0.00632	88.9762	3.593761	0.718795	0.02009
ZN	0.00000	100.0000	11.363636	13.615385	0.00000
INDUS	0.46000	27.7400	11.136779	7.078462	2.68000
CHAS	0.00000	1.0000	0.069170	0.153846	0.00000
NOX	0.38500	0.8710	0.554695	0.539238	0.41610
RM	3.56100	8.7800	6.284634	8.348538	8.03400
AGE	2.90000	100.0000	68.574901	71.538462	8.40000
DIS	1.12960	12.1265	3.795043	3.430192	1.80100
RAD	1.00000	24.0000	9.549407	7.461538	2.00000
TAX	187.00000	711.0000	408.237154	325.076923	224.00000
PTRATIO	12.60000	22.0000	18.455534	16.361538	13.00000
B	0.32000	396.9000	356.674032	385.210769	354.55000
LSTAT	1.73000	37.9700	12.653063	4.310000	2.47000
MDEV	5.00000	50.0000	22.532806	44.200000	21.90000

	max_RM>8
CRIM	3.47428
ZN	95.00000
INDUS	19.58000
CHAS	1.00000
NOX	0.71800
RM	8.78000
AGE	93.90000
DIS	8.90670
RAD	24.00000
TAX	666.00000
PTRATIO	20.20000
B	396.90000
LSTAT	7.44000
MDEV	50.00000

On average, the tracts with greater than 8 rooms have a ver low CRIM, lower than the overall CRIM mean, while the proportion of Black population is higher than the mean, closer to the maximum overall value. A thing to note is that they do have a high MDEV as well, meaning that these rooms are likely rented out.

4.a.

$Y = 50 + 20(\text{GPA}) + 0.07(\text{IQ}) + 35(\text{LVL}) + 0.01(\text{GPA} * \text{IQ}) - 10 (\text{GPA} * \text{LVL})$ college: (LVL = 1) $35 - 10 \text{ GPA}$ highschool: (LVL = 0) $0 (35 - 10 \text{ GPA}) - 0 = 35 - 10 \text{ GPA}$ GPA = 2, College students earn = $35 - 10 * 2 = 15\text{K}$ more

3, College students earn = $35 - 10 * 3 = 5\text{k}$ more

3.5, College students earn = $35 - 10 * 3.5 = \text{don't earn more}$

4, HS students = $35 - 10 * 4 = \text{earn } 5\text{k}$ more

The correct answer is: iii. For a fixed value of IQ and GPA, high school graduates earn more, on average, than college graduates provided that the GPA is high enough.

4.b. Predict the salary of a college graduate with IQ of 110 and a GPA of 4.0. $50 + 20(4) + 0.07(110) + 35(1) + 0.01(4 * 110) - 10 (4 * 1) = 50 + 80 + 7.7 + 35 + 4.4 - 40 = \137.1K

4.c. False. We need look at the standard error of the coefficient, as well as the sigma squared (unexplained variation in Y). We learned that the magnitudes and indiv hypothesis tests aren't really good ways to assess the models. Sometimes, a small magnitude might signify a big change (from GPA 3 to 4) vs IQ of 90 vs 130. Moreover, we want to look at it's relation to other predictors too.

5.a.

```

# List of indevependnt variables
independent_var = boston_df.columns[1:]

# Function to fit a model and collect results
def simple_reg(x):
    reg = smf.ols(f"CRIM ~ {x}", data=boston_df).fit()
    return {
        "predictor": x,
        "coefficient": reg.params.iloc[1],
        "p-value": reg.pvalues.iloc[1],
        "adjusted R-squared (single)": reg.rsquared_adj
    }

# Apply the function to each predictor and create a DataFrame
results_simple_reg = pd.DataFrame([simple_reg(x) for x in independent_var])

print(results_simple_reg)

```

	predictor	coefficient	p-value	adjusted R-squared (single)
0	ZN	-0.073521	6.151722e-06	0.037878
1	INDUS	0.506847	2.444137e-21	0.161937
2	CHAS	-1.871545	2.143436e-01	0.001080
3	NOX	30.975259	9.159490e-23	0.172686
4	RM	-2.691045	5.838094e-07	0.046485
5	AGE	0.107131	4.259064e-16	0.121309
6	DIS	-1.542831	1.268832e-18	0.141110
7	RAD	0.614137	1.620605e-55	0.385704
8	TAX	0.029563	9.759521e-47	0.334577
9	PTRATIO	1.144613	3.875122e-11	0.081269
10	B	-0.035535	1.432088e-18	0.140702
11	LSTAT	0.544406	7.124778e-27	0.202925
12	MDEV	-0.360647	2.083550e-19	0.147177

Sources: https://www.statsmodels.org/dev/generated/statsmodels.regression.linear_model.OLSResults.rsquared_adj.html
<https://stackoverflow.com/questions/41075098/how-to-get-the-p-value-in-a-variable-from-olsresults-in-python>

Plotting 5 predictors

```

# Loop
for x in independent_var[:5]:


```


```
fig, ax = plt.subplots(figsize=(8, 12))


# regression
sns.regplot(x=x, y='CRIM', data=boston_df, ax=ax)


ax.set_xlabel(x, fontsize=15)
ax.set_ylabel("CRIM", fontsize=15)


plt.show()
```











5.b.

```
mult_reg = smf.ols(  
    "CRIM ~ AGE + B + CHAS + DIS+ INDUS + NOX + RM + RAD + ZN + TAX + PTRATIO  
    ↪ + LSTAT + MDEV", data=boston_df).fit()  
  
print(mult_reg.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  CRIM    R-squared:
0.448
Model:                          OLS    Adj. R-squared:
0.434
Method:                         Least Squares    F-statistic:
30.73
Date:                           Thu, 23 Jan 2025    Prob (F-statistic):
2.04e-55
Time:                           23:16:30    Log-Likelihood:
-1655.7
No. Observations:                506    AIC:
3339.
Df Residuals:                    492    BIC:
3399.
Df Model:                        13
Covariance Type:                 nonrobust
=====
                                coef    std err          t      P>|t|      [0.025
                                0.975]
-----
Intercept      17.4184         7.270        2.396      0.017       3.135
31.702
AGE             0.0020         0.018        0.112      0.911      -0.033
0.037
B              -0.0069         0.004       -1.857      0.064      -0.014
0.000
CHAS           -0.7414         1.186       -0.625      0.532      -3.071
1.588
DIS            -0.9950         0.283       -3.514      0.000      -1.551
-0.439
INDUS          -0.0616         0.084       -0.735      0.463      -0.226
0.103
```

NOX	-10.6455	5.301	-2.008	0.045	-21.061
-0.230					
RM	0.3811	0.616	0.619	0.536	-0.829
1.591					
RAD	0.5888	0.088	6.656	0.000	0.415
0.763					
ZN	0.0449	0.019	2.386	0.017	0.008
0.082					
TAX	-0.0037	0.005	-0.723	0.470	-0.014
0.006					
PTRATIO	-0.2787	0.187	-1.488	0.137	-0.647
0.089					
LSTAT	0.1213	0.076	1.594	0.112	-0.028
0.271					
MDEV	-0.1992	0.061	-3.276	0.001	-0.319
-0.080					

```
=====
Omnibus:                662.271   Durbin-Watson:
1.515
Prob(Omnibus):          0.000   Jarque-Bera (JB):
82701.666
Skew:                   6.544   Prob(JB):
0.00
Kurtosis:               64.248   Cond. No.
1.58e+04
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.58e+04. This might indicate that there are strong multicollinearity or other numerical problems.

The coefficients represent the estimated effect of a one unit change in each of the independent variables, holding all other predictors constant. The intercept is the per capita crime rate when all the other predictors are 0. We can reject the null hypothesis at the 5% significance level for the ff predictors which have a p-value less than .05.: DIS, NOX, ZN, RAD, and MDEV for having low p-values.

5.c. Create a plot displaying the univariate regression coefficients from Question (5a) on the x-axis, and the multiple regression coefficients from Question (5b) on the y-axis.

```

# Univariate regression coefficients
coefs_df = pd.DataFrame({
    "predictor": results_simple_reg["predictor"],
    "uni_coefs": results_simple_reg["coefficient"]
})

# Add multivariate coefficients to the df
coefs_df["multi_coefs"] = mult_reg.params.loc[independent_var].values

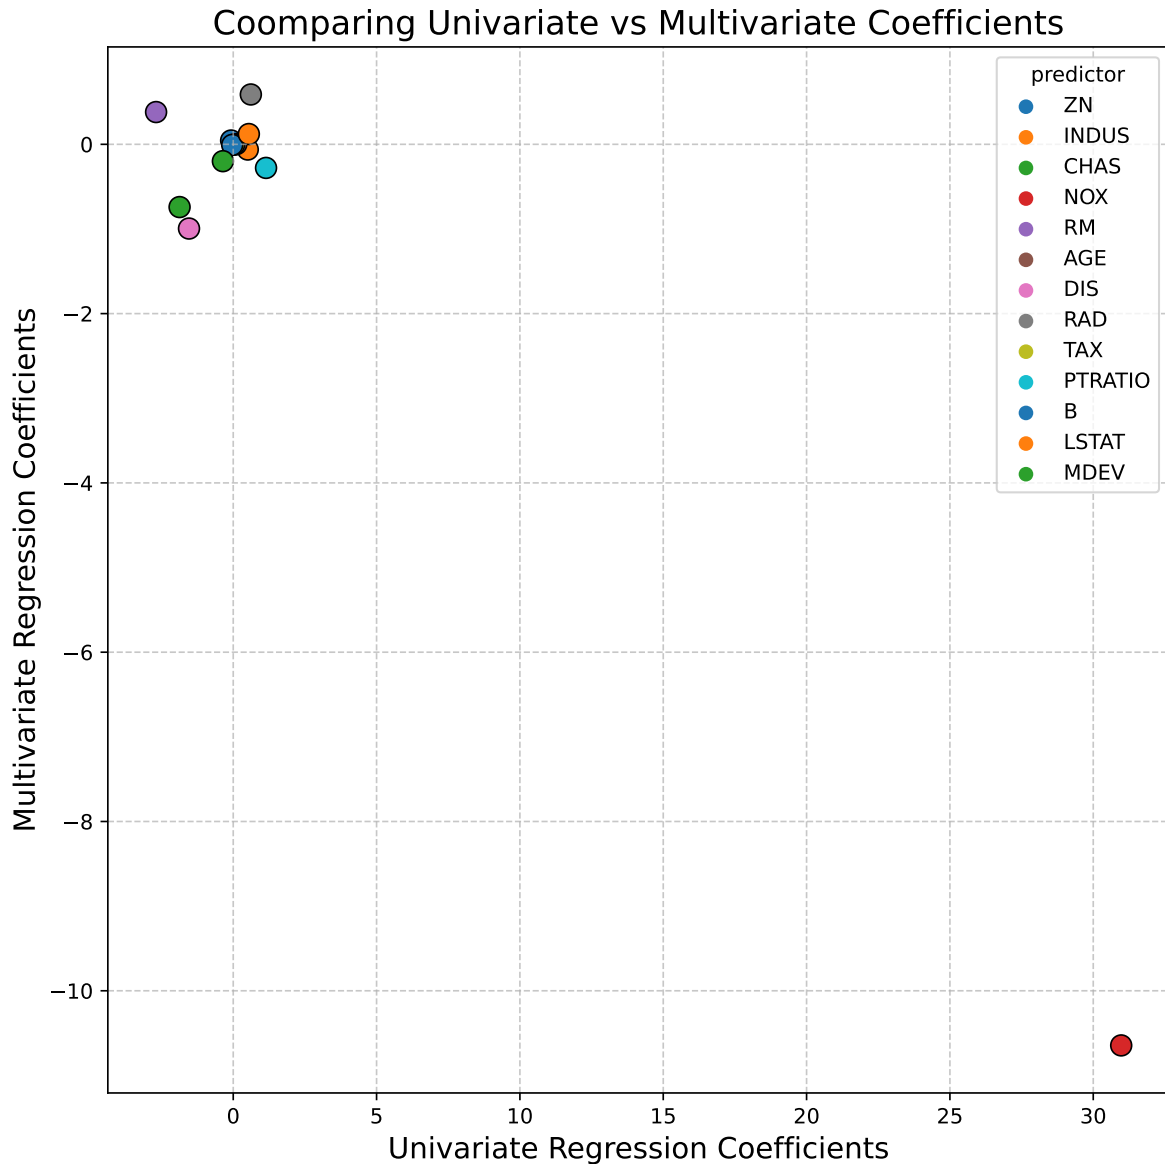
coefs_df = coefs_df.reset_index(drop=True)

# Plotting
fig, ax = plt.subplots(figsize=(8, 8))
sns.scatterplot(
    x="uni_coefs",
    y="multi_coefs",
    hue="predictor",
    palette="tab10",
    data=coefs_df,
    ax=ax,
    s=100,
    edgecolor="black"
)

# Set labels
ax.set_title(
    "Comparing Univariate vs Multivariate Coefficients", fontsize=16)
ax.set_xlabel("Univariate Regression Coefficients", fontsize=14)
ax.set_ylabel("Multivariate Regression Coefficients", fontsize=14)
ax.grid(True, linestyle="--", alpha=0.7)

# Show the plot
plt.tight_layout()
plt.show()

```



source:<https://stackoverflow.com/questions/42767489/add-legend-to-seaborn-point-plot>

*asked ChatGPT: Please help me refine with color palette an sizing.Helped me choose the best pallete and sizing to use for htis

Most of the data is clustered at the upperleft corner of the graphm wiith NOX being the one outlier (a high univariate coef value, but low mutivariate coef value). It looks liek the coefficients do not match. To observe the other data points more closely, we can redo the graph, without NOX so it will zoom in,

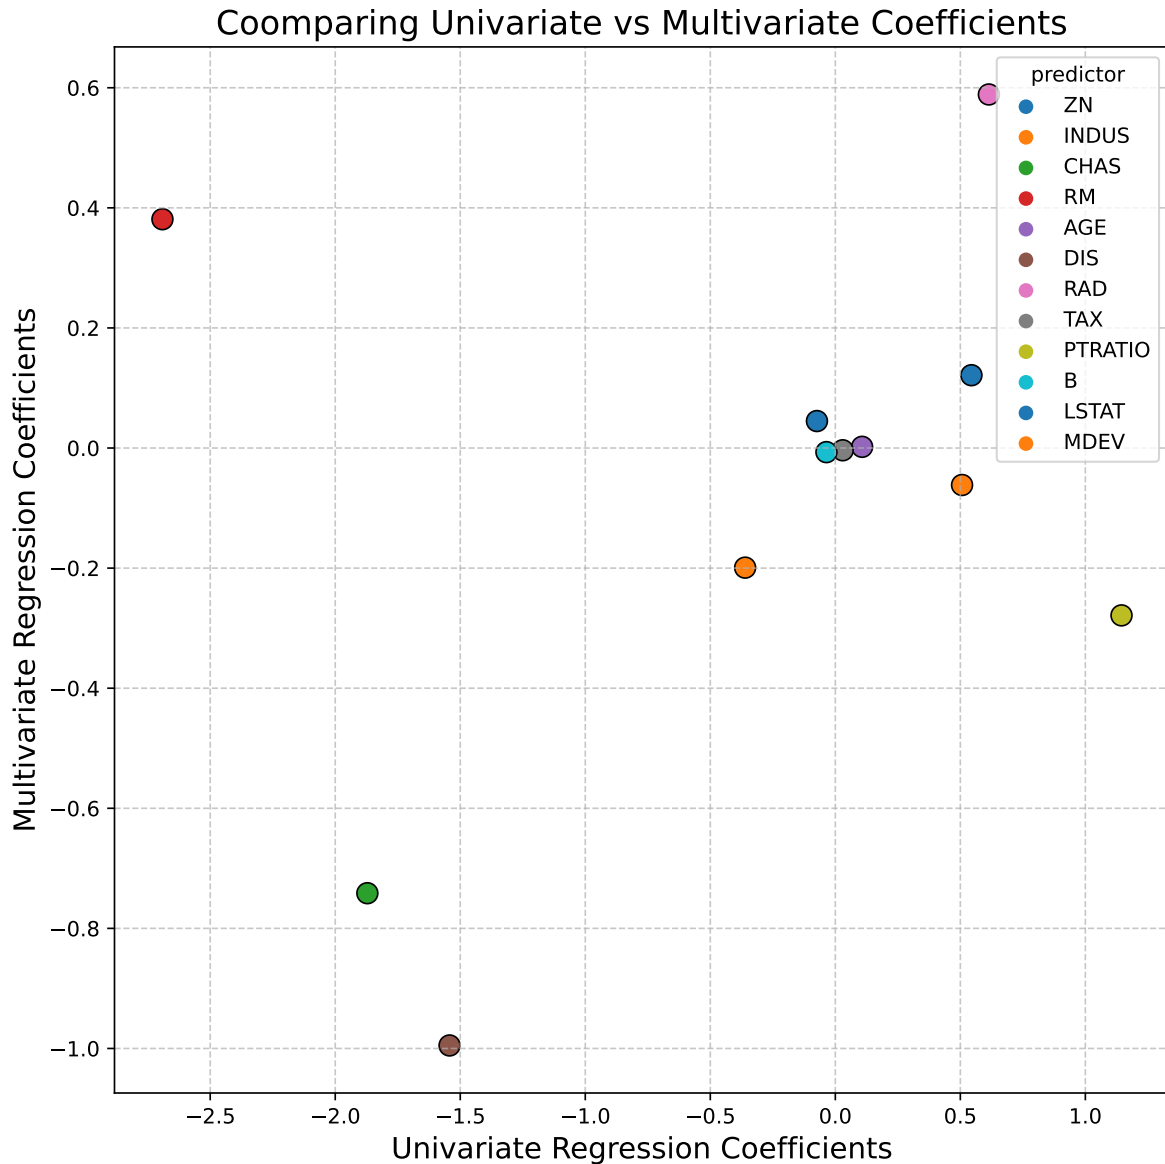

```

fig, ax = plt.subplots(figsize=(8, 8))
sns.scatterplot(
    x="uni_coefs",
    y="multi_coefs",
    hue="predictor",
    palette="tab10",
    data=coefs_df[coefs_df["predictor"]!="NOX"],
    ax=ax,
    s=100,
    edgecolor="black"
)

# Set labels
ax.set_title(
    "Comparing Univariate vs Multivariate Coefficients", fontsize=16)
ax.set_xlabel("Univariate Regression Coefficients", fontsize=14)
ax.set_ylabel("Multivariate Regression Coefficients", fontsize=14)
ax.grid(True, linestyle="--", alpha=0.7)

# Show the plot
plt.tight_layout()
plt.show()

```



Now, we can see that the data is actually spread out—there are both positive and negative values. Strangely, some values that are positive in the univariate regressio (PTRATIO and INDUS), are negative in the multivariate, and vice versa (RM and ZN).

5.d.

```
# Extracting adjusted R2 from earlier
linear_r2 = results_simple_reg[["predictor", "adjusted R-squared
↪ (single)"]].rename(
```

```

    columns={"adjusted R-squared (single)": "Adj R2 (Linear)"}
)

# Function for polynomial models and get adjusted R2
def fit_polynomial(data, predictor, response='CRIM'):
    # Fit polynomial model
    formula = f"{response} ~ {predictor} + I({predictor}**2) +
↪ I({predictor}**3)"
    model = smf.ols(formula, data=boston_df).fit()
    return model.rsquared_adj

# Fit polynomial models for each predictor
#empty list
polynomial_r2 = []
#loop and apply function
for predictor in linear_r2["predictor"]:
    adj_r2_poly = fit_polynomial(boston_df, predictor)
    polynomial_r2.append(
        {"predictor": predictor, "Adj R2 (Polynomial)": adj_r2_poly})

# Convert to DataFrame
polynomial_r2_df = pd.DataFrame(polynomial_r2)

# Merge results into a table
comparison_df = linear_r2.merge(polynomial_r2_df, on="predictor")
comparison_df["Difference"] = (
    comparison_df["Adj R2 (Linear)"] - comparison_df["Adj R2 (Polynomial)"]
)

from tabulate import tabulate
print(tabulate(comparison_df, headers="keys", tablefmt="grid"))

```

	predictor	Adj R2 (Linear)	Adj R2 (Polynomial)	Difference
0	ZN	0.0378783	0.0520163	-0.014138
1	INDUS	0.161937	0.252486	-0.0905491

+-----+-----+-----+-----+-----+				
	2	CHAS		0.00107951 -0.000910111 0.00198962
+-----+-----+-----+-----+-----+				
	3	NOX		0.172686 0.288145 -0.115459
+-----+-----+-----+-----+-----+				
	4	RM		0.0464854 0.0628415 -0.0163561
+-----+-----+-----+-----+-----+				
	5	AGE		0.121309 0.167472 -0.0461626
+-----+-----+-----+-----+-----+				
	6	DIS		0.14111 0.271521 -0.130411
+-----+-----+-----+-----+-----+				
	7	RAD		0.385704 0.391982 -0.00627877
+-----+-----+-----+-----+-----+				
	8	TAX		0.334577 0.361 -0.0264232
+-----+-----+-----+-----+-----+				
	9	PTRATIO		0.0812689 0.10718 -0.025911
+-----+-----+-----+-----+-----+				
	10	B		0.140702 0.13884 0.00186292
+-----+-----+-----+-----+-----+				
	11	LSTAT		0.202925 0.209716 -0.00679132
+-----+-----+-----+-----+-----+				
	12	MDEV		0.147177 0.412562 -0.265385
+-----+-----+-----+-----+-----+				

source: <https://www.geo.fu-berlin.de/en/v/soga-py/Basics-of-statistics/Linear-Regression/Polynomial-Regression/Polynomial-Regression—An-example/index.html>

table source: <https://www.datacamp.com/tutorial/python-tabulate>

For all non-indicator predictors, the adjusted R-squared is higher in the polynomial vs the simple regression, except for B. This means that the polynomial regression explains more of the variation in the model, thus fits the data better. It also means that the additional coomplexity in the model did generally add value.